

Project 1: Exploring Bandit Algorithms in Stationary and Non-Stationary Environments

Student: Harshani Rathnayake
Departments: Mathematics and Statistics and Computer Science
Date: June 15, 2025
GitHub Repository: <https://github.com/Harshani-Rathnayake/Project-1-Bandit-Algorithms.git>

Abstract

The multi-armed bandit (MAB) problem balances the competing objectives of exploration and exploitation by offering a straightforward framework for sequential decision-making under uncertainty. Several standard bandit algorithms, including the gradient bandit method, epsilon-greedy with tuned parameters, greedy algorithm, and optimistic initialization, are implemented and compared in this study under both stationary and non-stationary reward settings. The non-stationary environments are made to mimic the instability of reward systems in the actual world by incorporating mean-reverting dynamics, sudden changes in points, and smooth drift. Two main performance indicators are used to evaluate each approach over 1000 simulations: the percentage of optimal action selection and the average per-step reward. The results offer useful suggestions for algorithm selection in dynamic environments and empirical insight into the adaptability and learning capability of each method. The results and code are all made openly accessible and completely replicable.

Introduction

In artificial intelligence and reinforcement learning, sequential decision-making in unpredictable environments is a basic challenge. For investigating these problems, the multi-armed bandit (MAB) problem provides a standard framework. In its most basic form, an agent periodically chooses from a group of actions (or arms), each of which produces random rewards derived from unidentified distributions. The objective is to maximize cumulative reward over time, which calls for striking a balance between *exploitation*—using the existing knowledge to choose the most well-known action—and *exploration*—learning about lesser-known actions.[3]

In this study, several well-known multi-armed bandit algorithms are practically studied in various environmental settings. To be more precise, this paper compares:

- Greedy algorithms initialized with zero action values,
- Epsilon-greedy algorithms with exploration rates selected through pilot tuning,
- Optimistic initialization methods using domain knowledge,
- Gradient bandit algorithms with varying learning rates.

First, these algorithms are tested in a stationary environment with an initial reward distribution for each arm. The analysis is then extended to non-stationary environments, which are defined by either gradual drift, mean-reverting dynamics, or abrupt changes (changepoints). These hypothetical situations simulate real-life circumstances in which reward systems change over time as a result of shifting consumer habits, changes in the market, or environmental influences.[1]

Across 1000 independent simulations of 2000 time steps, the performance of each algorithm is based on the average per-step reward and the proportion of optimal actions selected. This research contributes to a better understanding of instructional techniques in practice by illuminating the relative advantages and disadvantages of each strategy in both stable and dynamic situations.

1 Part 1: Stationary Bandit Environment

This experiment uses the conventional k -armed bandit testbed with $k = 10$ to investigate and evaluate the performance of different classical bandit learning algorithms in a stationary environment. The true action-value $\mu_i \sim \mathcal{N}(0, 1)$ is calculated once every simulation and held constant in the reward distribution $\mathcal{N}(\mu_i, 1)$ that corresponds to each arm $i \in \{1, \dots, 10\}$. The learning algorithms are evaluated as follows:

- **Greedy (Q=0):** Pure exploitation with action-value estimates initialized at zero.
- **Epsilon-Greedy:** Selects a random action with probability ε , otherwise selects the greedy action. The best ε was tuned through pilot experiments.
- **Optimistic Greedy:** Greedy approach with optimistic initial values set to the 99.5th percentile of the best arm’s reward distribution to encourage exploration.
- **Gradient Bandit:** Policy-gradient method that uses preferences instead of value estimates and applies a softmax to select actions. The learning rate α was tuned experimentally.

Hyperparameter Tuning: Several choices of $\varepsilon \in \{0.01, 0.05, 0.1, 0.2\}$ were tested in pilot trials, and the value that produced the largest average reward in the last 100 steps across 200 simulations was selected. The gradient bandit technique was evaluated using a similar grid search strategy, examining $\alpha \in \{0.01, 0.05, 0.1, 0.4\}$. The 99.5th percentile of the best arm distribution averaged over several trials was the optimistic initial value.

Every algorithm was executed in 2000 time steps and replicated over 1000 independent simulations using various arm values and random seeds. The performance metrics are as follows:

1. Average reward at each time step.
2. Percentage of times the optimal action (arm with highest μ_i) was selected.

1.1 Results – Stationary Setting

The average reward per step for each method is compared in Figure 1. Lack of exploration causes the greedy algorithm to plateau too soon. Both gradient bandit and epsilon-greedy bandit techniques balance exploration and exploitation more effectively, resulting in larger average rewards.

The percentage of times the optimal action was selected over time is displayed in Figure 2, demonstrating how rapidly and reliably each approach learns to favour the better arm.

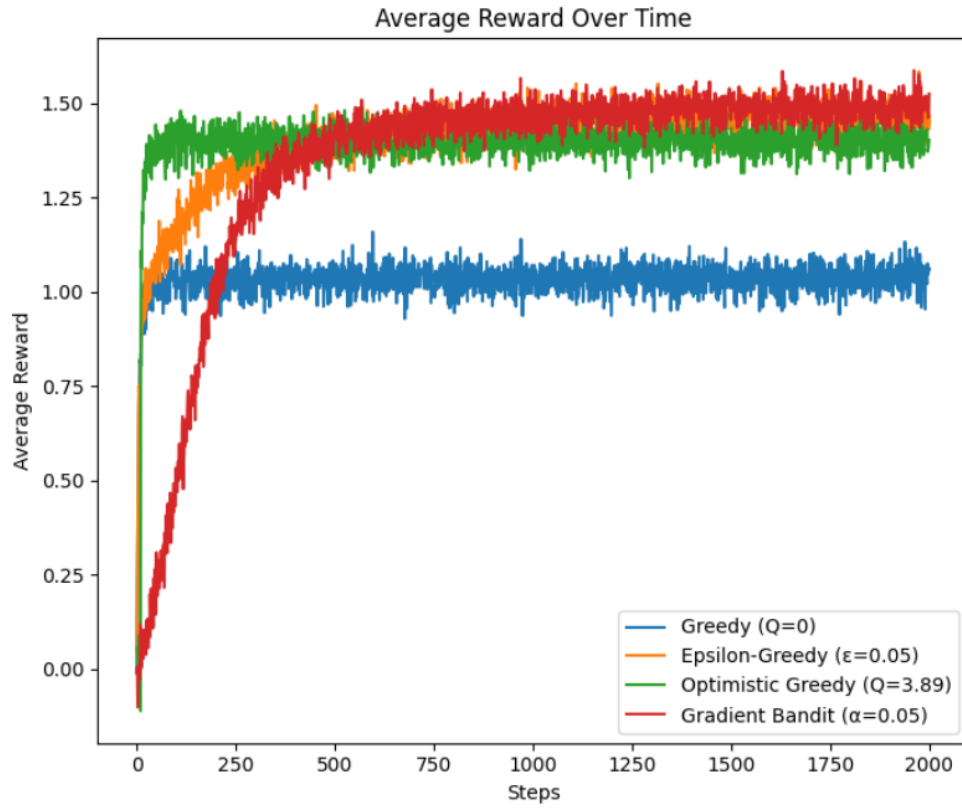


Figure 1: Average reward over 2000 steps for each algorithm.

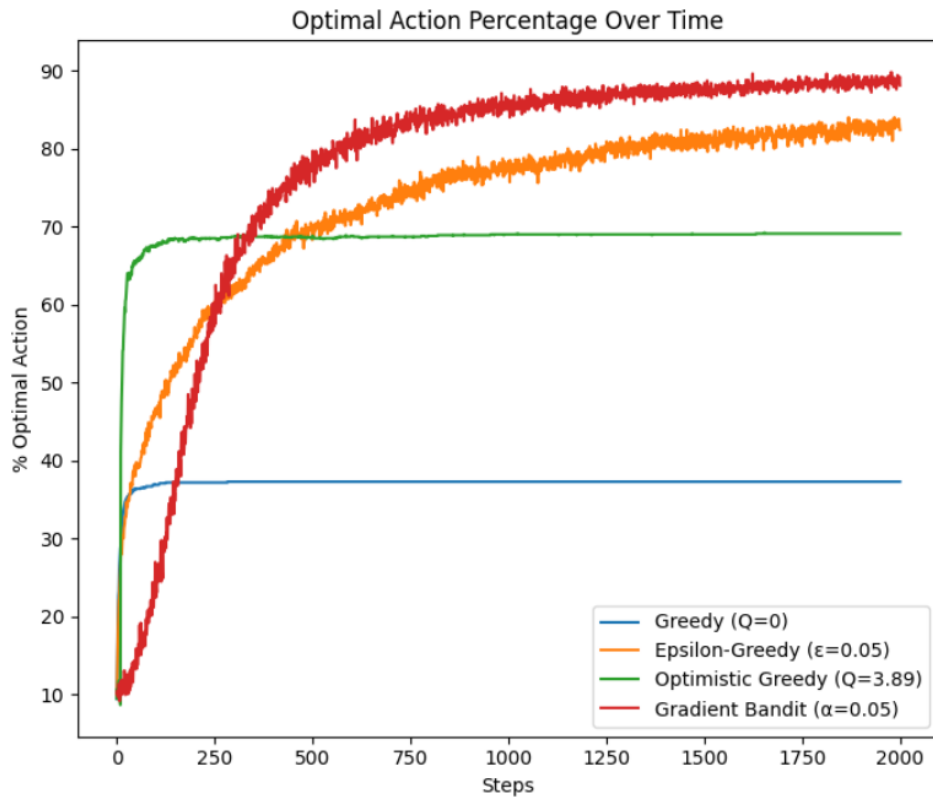


Figure 2: Percentage of time the optimal action is taken over 2000 steps for each algorithm.

Algorithm	Final Average Reward	Percentage Optimal Action
Greedy (Q=0)	1.043	37.30%
Epsilon-Greedy ($\varepsilon = 0.05$)	1.462	82.74%
Optimistic Greedy (Q=3.89)	1.399	69.10%
Gradient Bandit ($\alpha = 0.05$)	1.493	88.58%

Table 1: Performance metrics averaged over the final 100 steps.

1.2 Discussion – Part 1

The results make it extremely evident that exploration is essential to locating and taking advantage of the best action in stationary bandit environments. The low optimal action percentage and average reward show that the pure greedy algorithm, which lacks exploration, prematurely converges to ineffective actions.

The exploration and exploitation are successfully balanced by the epsilon-greedy strategy with $\varepsilon = 0.05$, which produces noticeably better results than greedy.

The optimistic greedy method promotes initial exploration with high starting values, but eventually behaves greedily, resulting in average performance.

Finally, the gradient bandit algorithm outperforms all others, likely due to its use of softmax action preferences, which allow it to continuously adapt based on reward feedback.

All things considered, the gradient bandit and epsilon-greedy algorithms perform better in stationary situations, with the gradient bandit obtaining the largest average reward and the best rate of action selection.

2 Part 2: Non-Stationary Bandit Environment

The reward distributions are frequently dynamic rather than constant in real-world situations. The multi-armed bandit problem is extended in this section to non-stationary environments, where the predicted rewards of actions change either suddenly or gradually. Under these circumstances, this study set several bandit algorithms into practice and evaluated them, assessing their adaptability by calculating average rewards and the percentage of times the most effective option of action is selected. [2]

2.1 Gradual Changes

The variations in the reward distributions over time replicate settings in which the underlying parameters gradually drift or return to a mean. Two categories of gradual changes were examined:

- **Drift:** The mean reward $\mu_{i,t}$ for each action is transformed into a random walk with Gaussian noise introduced, as modeled by

$$\mu_{i,t} = \mu_{i,t-1} + \epsilon_{i,t}, \quad \epsilon_{i,t} \sim \mathcal{N}(0, 0.0001)$$

- **Mean-Reverting:** The mean reward follows a mean-reverting process:

$$\mu_{i,t} = \kappa\mu_{i,t-1} + \epsilon_{i,t}, \quad \epsilon_{i,t} \sim \mathcal{N}(0, 0.0001)$$

with $\kappa = 0.5$, causing the means to fluctuate around zero with added noise.

These small changes simulate things like changing brightness over the course of a day or gradually shifting user preferences.

The average reward and optimal action percentage for each algorithm under the drift scenario are plotted side by side in the following Figure 3 below.

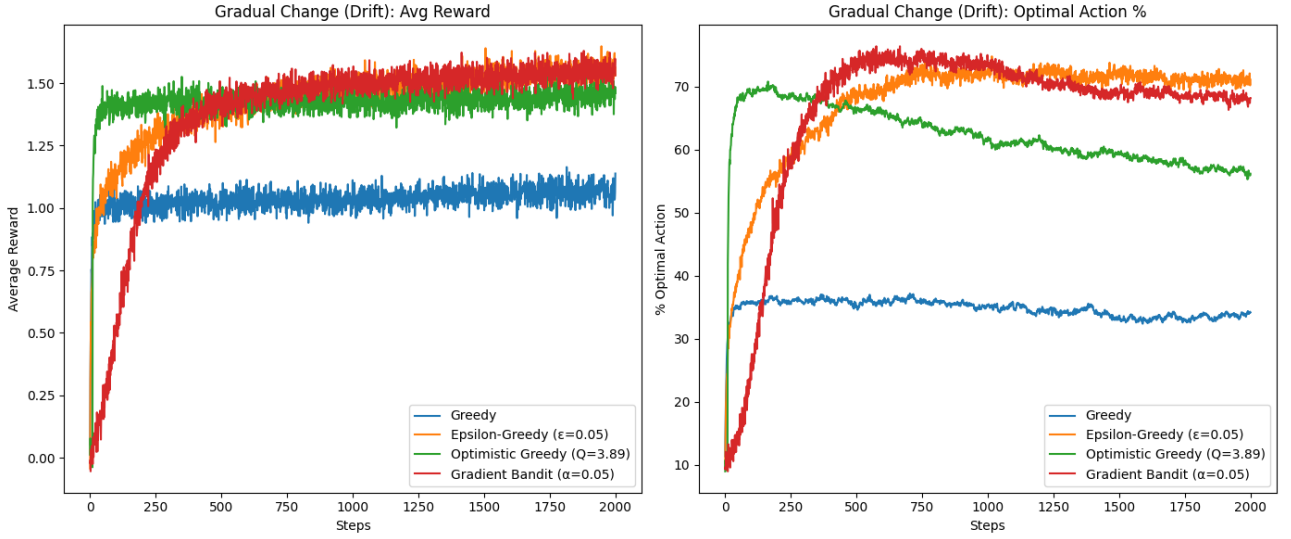


Figure 3: Gradual Change (Drift): Average Reward and Optimal Action Percentage for Different Bandit Algorithms

Comparable outcomes for the mean-reverting procedure are displayed in Figure 4 below.




Figure 4: Gradual Change (Mean-Reverting): Average Reward and Optimal Action Percentage for Different Bandit Algorithms

2.2 Abrupt changes

Abrupt changes simulate rapid shifts, like turning a switch, in reward distribution. The reward means are randomly switched between acts at step 501, signifying an unanticipated shift in the surroundings. The two examples were investigated:

- **No Reset:** The bandit algorithms use their update mechanisms to adapt, continuing without resetting their learned values.

- **With Reset:** To simulate a situation in which the agent is aware of the change and begins learning again from scratch, the algorithms are "hard reset" at step 501 by setting all action-value estimates back to zero.

The two sudden change scenarios are shown below side by side with average reward and optimal action percentage graphs in Figures 5 and 6.

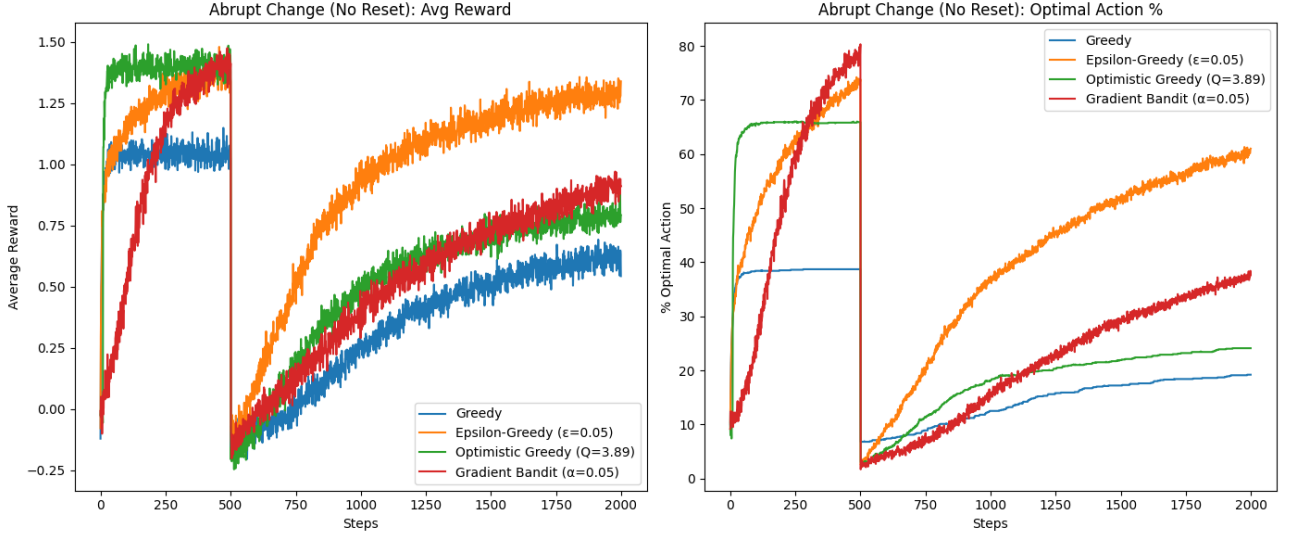


Figure 5: Abrupt Change (No Reset): Average Reward and Optimal Action Percentage for Different Bandit Algorithms

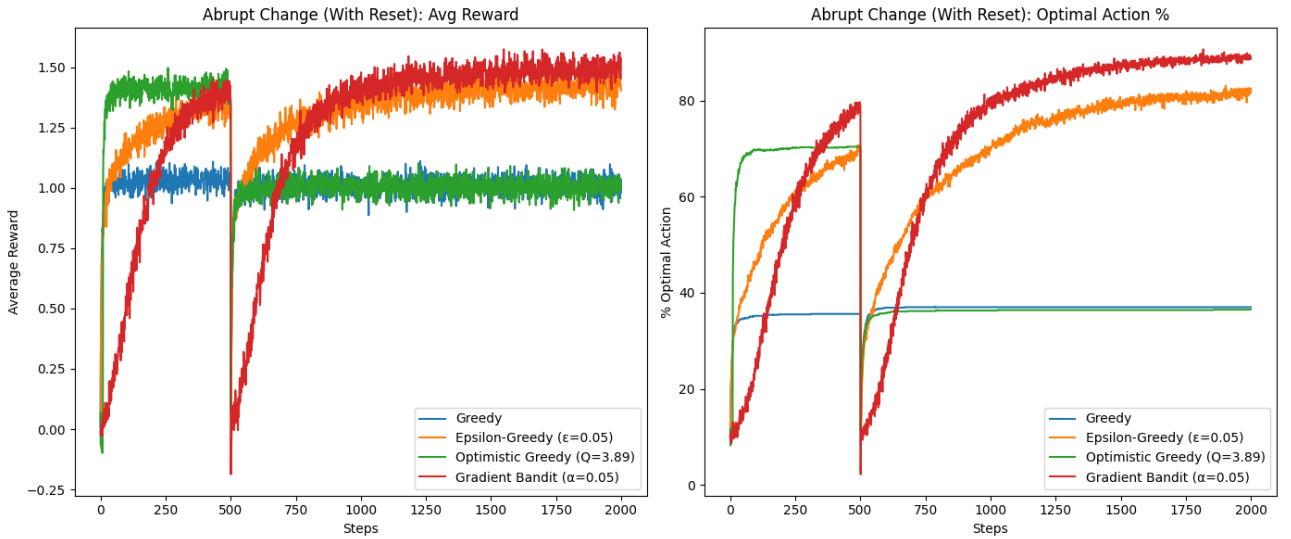


Figure 6: Abrupt Change (With Reset): Average Reward and Optimal Action Percentage for Different Bandit Algorithms

2.3 Results – Non-Stationary Setting

The performance reports for 2000 steps and 1000 runs provide numerous important insights:

Change Type	Algorithm	Average Reward	Percentage Optimal Action
Gradual Change (Drift)	Greedy	1.138	34.20%
	Epsilon-Greedy ($\epsilon = 0.05$)	1.563	71.10%
	Optimistic Greedy ($Q = 3.89$)	1.460	56.20%
	Gradient Bandit ($\alpha = 0.05$)	1.531	68.20%
Gradual Change (Mean-Reverting)	Greedy	0.028	9.60%
	Epsilon-Greedy ($\epsilon = 0.05$)	-0.035	9.90%
	Optimistic Greedy ($Q = 3.89$)	-0.029	8.60%
	Gradient Bandit ($\alpha = 0.05$)	-0.026	10.00%
Abrupt Change (No Reset)	Greedy	0.543	19.20%
	Epsilon-Greedy ($\epsilon = 0.05$)	1.339	61.00%
	Optimistic Greedy ($Q = 3.89$)	0.793	24.10%
	Gradient Bandit ($\alpha = 0.05$)	0.910	37.60%
Abrupt Change (With Reset)	Greedy	0.979	37.00%
	Epsilon-Greedy ($\epsilon = 0.05$)	1.402	82.50%
	Optimistic Greedy ($Q = 3.89$)	1.005	36.50%
	Gradient Bandit ($\alpha = 0.05$)	1.453	88.70%

Table 2: Final Summary of Bandit Algorithms Performance

- **Gradual Change (Drift):** The Epsilon-Greedy algorithm ($\epsilon = 0.05$) achieved the highest final average reward (1.563) and the highest percentage of optimal action selection (71.10%). The Gradient Bandit ($\alpha = 0.05$) also performed strongly, with a final average reward of 1.531 and 68.20% optimal actions. Greedy performed worst due to its lack of exploration, with a final average reward of 1.138 and only 34.20% optimal action selection.
- **Gradual Change (Mean-Reverting):** All methods struggled, showing near-zero or slightly negative average rewards and low optimal action percentages (around 8.6% to 10%). The mean-reverting nature of the environment likely caused rapid fluctuations that made stable learning difficult.
- **Abrupt Change (No Reset):** Performance decreased notably across all methods, especially Greedy and Optimistic Greedy, which showed lower average rewards (0.543 and 0.793 respectively) and optimal action percentages (19.20% and 24.10%). Epsilon-Greedy and Gradient Bandit adapted better but still suffered from the abrupt environmental changes, achieving 1.339 and 0.910 average rewards with 61.00% and 37.60% optimal actions respectively.
- **Abrupt Change (With Reset):** Explicitly resetting estimates at the change point improved adaptability significantly. Gradient Bandit ($\alpha = 0.05$) and Epsilon-Greedy ($\epsilon = 0.05$) achieved the highest final average rewards (1.453 and 1.402) and optimal action percentages (88.70% and 82.50%), substantially outperforming other methods.

2.4 Discussion – Part 2

These tests highlight the importance adaptability in non-stationary settings. Strong performance was continuously demonstrated by the Gradient Bandit algorithm, particularly in the drift and abrupt with reset scenarios. This is probably because the method is gradient-based and updates better monitor fluctuating reward structures. Although the epsilon-greedy technique did an outstanding job of balancing exploration, gradient methods somewhat overtook it.

All algorithms faced difficulties in the mean-reverting scenario, which highlighted the limitations when rewards oscillate unpredictably around a mean.

Reset or change-detection techniques are beneficial because abrupt changes without reset penalize algorithms that mostly count on accumulated past knowledge. In real-world applications, agents are capable to forget old information when their surroundings abruptly change.

All things considered, this section of the study demonstrated how several bandit algorithms handle time-varying reward distributions, highlighting the usefulness of integrating adaptation and reset techniques in sequential decision-making in the real world.

Conclusion

In conclusion, this study compared multiple bandit algorithms in stationary and non-stationary situations. The results revealed that exploration is critical, because the greedy approach performed poorly, whereas epsilon-greedy and gradient bandit strategies achieved bigger rewards and better action choices.

In dynamic circumstances, the gradient bandit algorithm generally performed well, particularly during steady drifts and rapid changes. Mean-reverting environments were problematic for all techniques, emphasizing the challenge of unstable rewards.

Moreover, reset learning after rapid changes boosted performance dramatically, demonstrating the need for adaptation. Overall, gradient-based and exploration-aware approaches are better adapted to real-world scenarios with dynamic reward structures.

Appendix

Part 1: Stationary Bandit Python Code

```
1 #Step 1: Bandit Environment
2 class Bandit:
3     def __init__(self, k=10):
4         #This initializes a k-armed bandit with true action values sampled
        from a normal distribution
5         self.k = k #number of arms
6         self.means = np.random.normal(0, 1, k) #true reward means for each
        arm
7         self.optimal_action = np.argmax(self.means) #index of the best arm
        (highest mean)
8     def step(self, action):
9         #This generates a stochastic reward for the selected action (arm)
10        return np.random.normal(self.means[action], 1) # reward is drawn
        from N(true_mean, 1)
11
12 def argmax_random_tie(values):
13     #This returns the index of the maximum value in the array, breaking ties
        randomly
14     max_val = np.max(values)
15     candidates = np.flatnonzero(values == max_val) #indices with max value
16     return np.random.choice(candidates)
17
18
19 #Step 2: Action Selection Algorithms
20 def greedy(bandit, steps=2000):
21     #Greedy strategy: always pick the current best-known action (Q)
22     Q = np.zeros(bandit.k) #estimated rewards
23     N = np.zeros(bandit.k) #count of times each arm was selected
24     rewards = []
25     optimal_actions = []
26     for t in range(steps):
27         action = argmax_random_tie(Q) #select best current action
28         reward = bandit.step(action)
29         N[action] += 1
30         Q[action] += (reward - Q[action]) / N[action] #incremental update
31         rewards.append(reward)
32         optimal_actions.append(int(action == bandit.optimal_action))
33     return rewards, optimal_actions
34
35 def epsilon_greedy(bandit, epsilon=0.1, steps=2000):
36     #Epsilon-greedy strategy: with probability epsilon, explore; otherwise
        exploit
37     Q = np.zeros(bandit.k)
38     N = np.zeros(bandit.k)
39     rewards = []
40     optimal_actions = []
41     for t in range(steps):
42         if np.random.rand() < epsilon:
43             action = np.random.randint(bandit.k) #explore
44         else:
45             action = argmax_random_tie(Q) #exploit
46         reward = bandit.step(action)
47         N[action] += 1
48         Q[action] += (reward - Q[action]) / N[action]
49         rewards.append(reward)
50         optimal_actions.append(int(action == bandit.optimal_action))
```

```

51     return rewards, optimal_actions
52
53 def optimistic_greedy(bandit, initial_value=5.0, steps=2000):
54     #Optimistic greedy: initializes estimates with high values to encourage
    initial exploration
55     Q = np.ones(bandit.k) * initial_value #optimistic start
56     N = np.zeros(bandit.k)
57     rewards = []
58     optimal_actions = []
59     for t in range(steps):
60         action = argmax_random_tie(Q)
61         reward = bandit.step(action)
62         N[action] += 1
63         Q[action] += (reward - Q[action]) / N[action]
64         rewards.append(reward)
65         optimal_actions.append(int(action == bandit.optimal_action))
66     return rewards, optimal_actions
67
68 def gradient_bandit(bandit, alpha=0.1, steps=2000):
69     #Gradient bandit algorithm using preference values and softmax action
    selection
70     H = np.zeros(bandit.k) #preferences
71     avg_reward = 0
72     rewards = []
73     optimal_actions = []
74     for t in range(steps):
75         probs = np.exp(H) / np.sum(np.exp(H)) #softmax over preferences
76         action = np.random.choice(bandit.k, p=probs) #pick action based on
    probabilities
77         reward = bandit.step(action)
78         avg_reward += (reward - avg_reward) / (t + 1) #update average
    reward
79         #Update preferences:
80         for a in range(bandit.k):
81             H[a] += alpha * ((1 if a == action else 0) - probs[a]) * (reward
    - avg_reward)
82         rewards.append(reward)
83         optimal_actions.append(int(action == bandit.optimal_action))
84     return rewards, optimal_actions
85
86 #Step 3: Run a Simulation
87 def run_simulation(algo_func, runs=1000, **kwargs):
88     #Below runs a bandit algorithm multiple times and average the results
89     steps = kwargs.get('steps', 2000)
90     avg_rewards = np.zeros(steps)
91     optimal_action_counts = np.zeros(steps)
92     for _ in range(runs):
93         bandit = Bandit()
94         rewards, optimal_actions = algo_func(bandit, **kwargs)
95         avg_rewards += rewards
96         optimal_action_counts += optimal_actions
97     avg_rewards /= runs
98     optimal_action_counts = (optimal_action_counts / runs) * 100 #convert
    to percentage
99     return avg_rewards, optimal_action_counts
100
101 #Step 4: Tune the Initial Value for Optimistic Greedy
102 optimistic_values = []
103 for _ in range(20):
104     b = Bandit()

```

```

105     max_mean = max(b.means)
106     # Calculate a high quantile value from a normal distribution centered on
107     # max_mean
108     percentile = np.percentile(np.random.normal(loc=max_mean, scale=1, size
109     =100000), 99.5)
110     optimistic_values.append(percentile)
111     #Below is the final average optimistic Q-value
112     optimistic_value = np.mean(optimistic_values)
113
114 #Step 5: Tune Epsilon
115 epsilons = [0.01, 0.05, 0.1, 0.2]
116 epsilon_results = {}
117 #Trying different epsilons and record final average rewards
118 for eps in epsilons:
119     rewards, _ = run_simulation(epsilon_greedy, epsilon=eps, runs=200, steps
120     =500)
121     epsilon_results[eps] = np.mean(rewards[-100:]) #average of last 100
122     steps
123
124 #Choose epsilon that gave highest average reward
125 best_eps = max(epsilon_results, key=epsilon_results.get)
126
127 #Step 6: Tune Alpha
128 alphas = [0.01, 0.05, 0.1, 0.4]
129 alpha_results = {}
130 #Tring different alphas and record final average rewards
131 for a in alphas:
132     rewards, _ = run_simulation(gradient_bandit, alpha=a, runs=200, steps
133     =500)
134     alpha_results[a] = np.mean(rewards[-100:])
135
136 #Choose the best alpha value
137 best_alpha = max(alpha_results, key=alpha_results.get)
138
139 #Step 7: Final Run of the Algorithms
140 results = {
141     "Greedy (Q=0)": run_simulation(greedy),
142     f"Epsilon-Greedy ( $\epsilon$ ={best_eps})": run_simulation(epsilon_greedy, epsilon
143     =best_eps),
144     f"Optimistic Greedy (Q={optimistic_value:.2f})": run_simulation(
145     optimistic_greedy, initial_value=optimistic_value),
146     f"Gradient Bandit ( $\alpha$ ={best_alpha})": run_simulation(gradient_bandit,
147     alpha=best_alpha)
148 }
149
150 #Step 8: Ploting the Results
151 def plot_results(results):
152     plt.figure(figsize=(14, 6))
153     #Plot average reward:
154     plt.subplot(1, 2, 1)
155     for label, (rewards, _) in results.items():
156         plt.plot(rewards, label=label)
157     plt.xlabel("Steps")
158     plt.ylabel("Average Reward")
159     plt.title("Average Reward Over Time")
160     plt.legend()
161
162     #Plot optimal action percentage:

```

```

157 plt.subplot(1, 2, 2)
158 for label, (_, optimal_actions) in results.items():
159     plt.plot(optimal_actions, label=label)
160 plt.xlabel("Steps")
161 plt.ylabel("% Optimal Action")
162 plt.title("Optimal Action Percentage Over Time")
163 plt.legend()
164 plt.tight_layout()
165 plt.show()
166 #Calling the plot function
167 plot_results(results)
168
169 #Step 9: Final Summary Output
170 #Below prints the final performance for each algorithm
171 for label, (rewards, optimal) in results.items():
172     final_avg_reward = np.mean(rewards[-100:]) #last 100 steps
173     final_opt_action = np.mean(optimal[-100:]) #last 100 steps
174     print(f"{label}: Final Avg Reward = {final_avg_reward:.3f}, Final %
Optimal = {final_opt_action:.2f}%")

```

Listing 1: Python Code for Stationary Bandit Algorithms

Part 2: Non-Stationary Bandit Python Code

```

1 #Non-Stationary Bandit Class:
2 class NonStationaryBandit:
3     def __init__(self, k=10, change_type=None, seed=None, permutation=None,
4         noise_stream=None):
5         self.k = k #number of arms (actions)
6         rng = np.random.default_rng(seed) #reproducible random generator
7         with optional seed
8         self.means = rng.normal(0, 1, k) #initialize mean reward for each
9         arm from N(0,1)
10        self.original_means = self.means.copy() #backup of original means (
11        for comparison or reset)
12        self.optimal_action = np.argmax(self.means) #initially optimal
13        action (arm with highest mean)
14        self.change_type = change_type #type of environment change ('drift
15        ', 'abrupt', etc)
16        self.permutation = permutation #used in abrupt change (permutes the
17        arm means)
18        self.noise_stream = noise_stream #stream of noise vectors for
19        gradual changes
20
21    def step(self, t, action):
22        #Apply changes to the means based on the type of environment
23        if self.change_type == 'drift':
24            #Gradually drift the means with small noise
25            if t > 0:
26                self.means += self.noise_stream[t]
27        elif self.change_type == 'mean_reverting':
28            #Move means closer to zero with some noise
29            if t > 0:
30                self.means = 0.5 * self.means + self.noise_stream[t]
31        elif self.change_type == 'abrupt' and t == 501:
32            #At time t=501, apply abrupt permutation of means
33            self.means = self.means[self.permutation]
34        #Update optimal action after changes
35        self.optimal_action = np.argmax(self.means)
36        #Return a sample reward from N(mean[action], 1)

```

```

29         return np.random.normal(self.means[action], 1)
30
31 #Runner for Non-Stationary Bandit Algorithms:
32 def run_nonstationary(algo_func, change_type, reset=False, **kwargs):
33     runs = kwargs.get('runs', 1000)
34     steps = kwargs.get('steps', 2000)
35     avg_rewards = np.zeros(steps) #to accumulate average rewards across all
    runs
36     optimal_action_counts = np.zeros(steps) #to accumulate % optimal
    actions
37     seeds = np.arange(runs)
38     #Create a fixed permutation for abrupt changes
39     permutation = np.random.permutation(10)
40     #Create noise streams for each run - used in drift or mean-reverting
41     noise_streams = [np.random.normal(0, 0.012, (steps, 10)) for _ in range(
    runs)]
42     #Run simulation over multiple independent seeds
43     for seed in seeds:
44         bandit = NonStationaryBandit(
45             change_type=change_type,
46             seed=seed,
47             permutation=permutation,
48             noise_stream=noise_streams[seed]
49         )
50         #Initialization of parameters
51         Q = np.zeros(bandit.k) #estimated value of each arm
52         N = np.zeros(bandit.k) #number of times each arm was selected
53         H = np.zeros(bandit.k) #preference values for gradient bandit
54         avg_reward = 0 #running average of reward for gradient bandit
55         #Extract tuning hyperparameters
56         alpha = kwargs.get('alpha', 0.1)
57         epsilon = kwargs.get('epsilon', 0.1)
58         init_val = kwargs.get('initial_value', 0)
59         #For optimistic method, set high initial Q values
60         if algo_func.__name__ == "optimistic_greedy":
61             Q = np.ones(bandit.k) * init_val
62         rewards = []
63         optimal_actions = []
64         for t in range(steps):
65             #Action selection based on algorithm type
66             if algo_func.__name__ == "epsilon_greedy":
67                 if np.random.rand() < epsilon:
68                     action = np.random.randint(bandit.k)
69                 else:
70                     action = argmax_random_tie(Q)
71             elif algo_func.__name__ in ["greedy", "optimistic_greedy"]:
72                 action = argmax_random_tie(Q)
73             elif algo_func.__name__ == "gradient_bandit":
74                 probs = np.exp(H) / np.sum(np.exp(H))
75                 action = np.random.choice(bandit.k, p=probs)
76             #Get reward from bandit
77             reward = bandit.step(t, action)
78             #Optional reset at time of abrupt change
79             if reset and t == 501:
80                 Q = np.zeros(bandit.k)
81                 N = np.zeros(bandit.k)
82                 H = np.zeros(bandit.k)
83                 avg_reward = 0
84             #Update estimates or preferences
85             if algo_func.__name__ in ["greedy", "epsilon_greedy", "

```

```

optimistic_greedy"]:
86         N[action] += 1
87         Q[action] += (reward - Q[action]) / N[action]
88         elif algo_func.__name__ == "gradient_bandit":
89             avg_reward += (reward - avg_reward) / (t + 1)
90             for a in range(bandit.k):
91                 H[a] += alpha * ((1 if a == action else 0) - probs[a]) *
(reward - avg_reward)
92             #Record metrics
93             rewards.append(reward)
94             optimal_actions.append(int(action == bandit.optimal_action))
95             avg_rewards += rewards
96             optimal_action_counts += optimal_actions
97         #Average metrics across runs
98         avg_rewards /= runs
99         optimal_action_counts = (optimal_action_counts / runs) * 100
100         return avg_rewards, optimal_action_counts
101
102
103 #Plotting Function for Non-Stationary Experiments:
104 def plot_nonstationary(results, title):
105     plt.figure(figsize=(14, 6))
106     #Plot: Average reward over Time
107     plt.subplot(1, 2, 1)
108     for label, (rewards, _) in results.items():
109         plt.plot(rewards, label=label)
110     plt.xlabel("Steps")
111     plt.ylabel("Average Reward")
112     plt.title(f"{title}: Avg Reward")
113     plt.legend()
114     #Plot: Optimal action percentage over Time
115     plt.subplot(1, 2, 2)
116     for label, (_, optimal_actions) in results.items():
117         plt.plot(optimal_actions, label=label)
118     plt.xlabel("Steps")
119     plt.ylabel("% Optimal Action")
120     plt.title(f"{title}: Optimal Action %")
121     plt.legend()
122     plt.tight_layout()
123     plt.show()
124
125 #Run Experiments for Different Change Types:
126 #Gradual Drift Change (non-stationary rewards increase/decrease slowly):
127 drift_results = {
128     "Greedy": run_nonstationary(greedy, change_type='drift'),
129     f"Epsilon-Greedy ( $\epsilon$ ={{best_eps}})": run_nonstationary(epsilon_greedy,
change_type='drift', epsilon=best_eps),
130     f"Optimistic Greedy (Q={{optimistic_value:.2f}})": run_nonstationary(
optimistic_greedy, change_type='drift', initial_value=optimistic_value),
131     f"Gradient Bandit ( $\alpha$ ={{best_alpha}})": run_nonstationary(gradient_bandit,
change_type='drift', alpha=best_alpha)
132 }
133 plot_nonstationary(drift_results, "Gradual Change (Drift)")
134
135 #Gradual Mean-Reverting Change (rewards decay towards zero):
136 reverting_results = {
137     "Greedy": run_nonstationary(greedy, change_type='mean_reverting'),
138     f"Epsilon-Greedy ( $\epsilon$ ={{best_eps}})": run_nonstationary(epsilon_greedy,
change_type='mean_reverting', epsilon=best_eps),
139     f"Optimistic Greedy (Q={{optimistic_value:.2f}})": run_nonstationary(

```

```

    optimistic_greedy, change_type='mean_reverting', initial_value=
    optimistic_value),
140     f"Gradient Bandit ( $\alpha$ ={{best_alpha}})": run_nonstationary(gradient_bandit,
    change_type='mean_reverting', alpha=best_alpha)
141 }
142 plot_nonstationary(reverting_results, "Gradual Change (Mean-Reverting)")
143
144 #Abrupt Change WITHOUT Reset (agent continues using old Q/H):
145 abrupt_norestart_results = {
146     "Greedy": run_nonstationary(greedy, change_type='abrupt', reset=False),
147     f"Epsilon-Greedy ( $\epsilon$ ={{best_eps}})": run_nonstationary(epsilon_greedy,
    change_type='abrupt', epsilon=best_eps, reset=False),
148     f"Optimistic Greedy (Q={{optimistic_value:.2f}})": run_nonstationary(
    optimistic_greedy, change_type='abrupt', initial_value=optimistic_value,
    reset=False),
149     f"Gradient Bandit ( $\alpha$ ={{best_alpha}})": run_nonstationary(gradient_bandit,
    change_type='abrupt', alpha=best_alpha, reset=False)
150 }
151 plot_nonstationary(abrupt_norestart_results, "Abrupt Change (No Reset)")
152
153 #Abrupt Change WITH Reset (agent resets internal estimates at t=501):
154 abrupt_reset_results = {
155     "Greedy": run_nonstationary(greedy, change_type='abrupt', reset=True),
156     f"Epsilon-Greedy ( $\epsilon$ ={{best_eps}})": run_nonstationary(epsilon_greedy,
    change_type='abrupt', epsilon=best_eps, reset=True),
157     f"Optimistic Greedy (Q={{optimistic_value:.2f}})": run_nonstationary(
    optimistic_greedy, change_type='abrupt', initial_value=optimistic_value,
    reset=True),
158     f"Gradient Bandit ( $\alpha$ ={{best_alpha}})": run_nonstationary(gradient_bandit,
    change_type='abrupt', alpha=best_alpha, reset=True)
159 }
160 plot_nonstationary(abrupt_reset_results, "Abrupt Change (With Reset)")
161
162 #Final Summary of the values:
163 def print_final_summary(results, title):
164     print(f"\n=== Final Summary for {title} ===")
165     for label, (avg_rewards, optimal_actions) in results.items():
166         final_avg_reward = avg_rewards[-1]
167         final_optimal_pct = optimal_actions[-1]
168         print(f"{{label}}: Final Avg Reward = {{final_avg_reward:.3f}}, Final %
    Optimal Action = {{final_optimal_pct:.2f}}%")
169     print("="*40)
170 #Print summary of each experiment:
171 print_final_summary(drift_results, "Gradual Change (Drift)")
172 print_final_summary(reverting_results, "Gradual Change (Mean-Reverting)")
173 print_final_summary(abrupt_norestart_results, "Abrupt Change (No Reset)")
174 print_final_summary(abrupt_reset_results, "Abrupt Change (With Reset)")

```

Listing 2: Python Code for Non-Stationary Bandit Algorithms

References

- [1] Peter Auer, Nicolò Cesa-Bianchi, and Paul Fischer. Finite-time analysis of the multiarmed bandit problem. *Machine learning*, 47(2):235–256, 2002.
- [2] Aurélien Garivier and Éric Moulines. Non-stationary bandit algorithms. In *Advances in Neural Information Processing Systems*, pages 1769–1776, 2008.
- [3] Richard S. Sutton and Andrew G. Barto. *Reinforcement Learning: An Introduction*. MIT Press, 2nd edition, 2018.